

Tuning Parameters

(based on tmwr.org)

Matthew McDonald

Hyperparamters

- Some parameters required for prediction can be estimated directly from the training data,
- Other parameters, called *tuning parameters* or *hyperparameters*, must be specified ahead of time and can't be directly found from training data.
- These are unknown structural or other kind of values that have significant impact on the model.

OLS Model Parameters

In ordinary linear regression, there are two parameters β_0 and β_1 of the model:

$$y_i = \beta_0 + \beta_1 x_i + \epsilon_i$$

When we have the outcome (y) and predictor (x) data, we can estimate the two parameters β_0 and β_1 :

$$\hat{\beta}_1 = \frac{\sum_i (y_i - \bar{y})(x_i - \bar{x})}{\sum_i (x_i - \bar{x})^2}$$

and

$$\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}.$$

K-Nearest Neighbor Model

K-nearest neighbors stores the training set (including the outcome).

When a new sample is predicted, K training set points are found that are most similar to the new sample being predicted.

The predicted value for the new sample is some summary statistic of the neighbors, usually:

- the mean for regression, or
- the mode for classification.

KNN Model Parameters

For the KNN model, the prediction equation for a new value x_0 is

$$\hat{y} = \frac{1}{K} \sum_{\ell=1}^K x_{\ell}^*$$

- K is the number of neighbors and the x_{ℓ}^* are the K closest values to x_0 in the training set.
- The model itself is not defined by a model equation
- The number of neighbors has a profound impact on the model; it governs the flexibility of the class boundary.
- For small values of K , the boundary is very elaborate while for large values, it might be quite smooth.

Note on KNN Model

Since the model is measuring distance, we typically should add a pre-processing step to center and scale all numeric parameters to ensure they're on the same scale.

Fitting a KNN Model to Ames

```
1 knn_mod <-
2   nearest_neighbor(neighbors = 5) %>%
3   set_engine("kkn") %>%
4   set_mode("regression")
5
6 # since Longitude and Latitude are already on the same scale,
7 # we can get away without centering and scaling
8 knn_wflow <-
9   workflow() %>%
10    add_formula(Sale_Price ~ Longitude + Latitude) %>%
11    add_model(knn_mod)
12
13 set.seed(1001)
14 ames_folds <- vfold_cv(ames_train, v = 10)
15
16 knn_fit <- knn_wflow %>% fit_resamples(resamples = ames_folds)
17 collect_metrics(knn_fit)
```

A tibble: 2 × 6

	.metric	.estimator	mean	n	std_err	.config
	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	rmse	standard	0.0987	10	0.00352	Preprocessor1_Model11
2	rsq	standard	0.686	10	0.0179	Preprocessor1_Model11

Setting K = 100

```
1 knn_mod <-
2   nearest_neighbor(neighbors = 100) %>%
3   set_engine("kknn") %>%
4   set_mode("regression")
5
6 knn_wflow <-
7   knn_wflow %>%
8   remove_model() %>%
9   add_model(knn_mod)
10
11 knn_fit <- knn_wflow %>% fit_resamples(resamples = ames_folds)
12 collect_metrics(knn_fit)
```

```
# A tibble: 2 × 6
  .metric .estimator mean      n std_err .config
  <chr>   <chr>      <dbl> <int>   <dbl> <chr>
1 rmse    standard    0.115     10 0.00335 Preprocessor1_Model11
2 rsq     standard    0.574     10 0.0151  Preprocessor1_Model11
```


What is the Best Choice for K ?

The tune() Function

How can we signal to tidymodels functions which arguments should be optimized? Parameters are marked for tuning by assigning them a value of `tune()`.

The `tune()` function doesn't execute any particular parameter value; it only returns an expression:

```
1 tune()
```

```
tune()
```

Embedding this `tune()` value in an argument will tag the parameter for optimization.

Tuning our KNN Model

```
1 knn_mod <-  
2   nearest_neighbor(neighbors = tune('K')) %>%  
3   set_engine("kkn") %>%  
4   set_mode("regression")  
5  
6 knn_wflow <-  
7   knn_wflow %>%  
8   remove_model() %>%  
9   add_model(knn_mod)  
10  
11 knn_tune <- knn_wflow %>% tune_grid(resamples = ames_folds,  
12                                     grid = tibble(K=1:20),  
13                                     metrics=metric_set(rmse))  
14 collect_metrics(knn_tune)
```

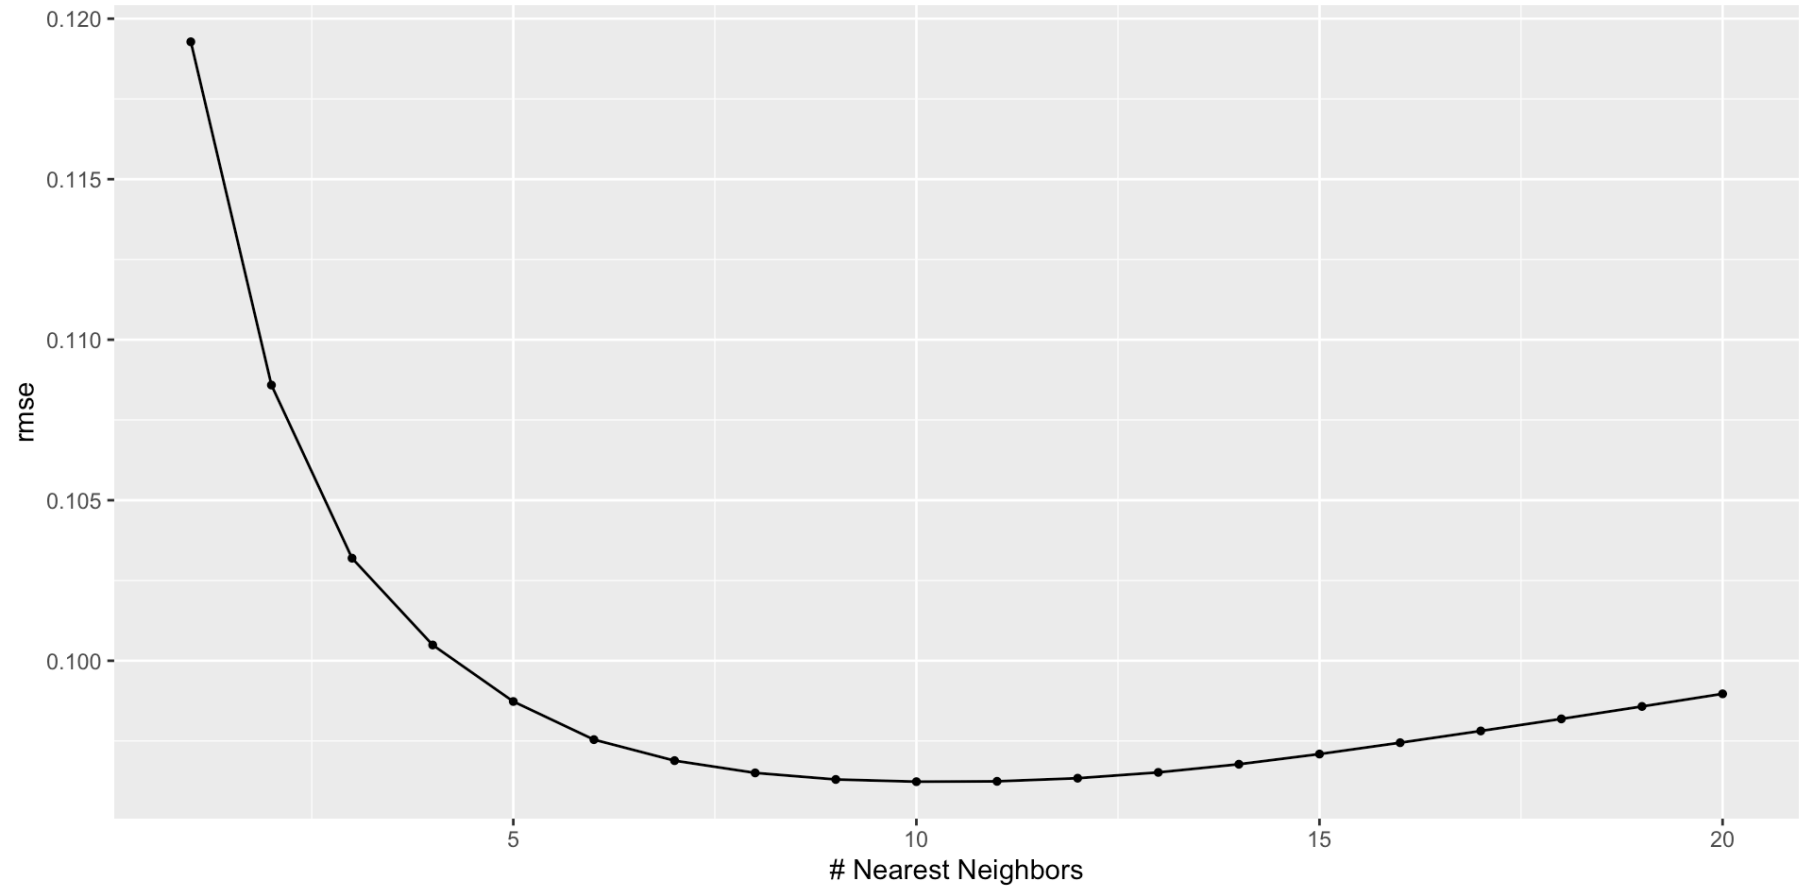
Tuning our KNN Model

A tibble: 20 × 7

	K	.metric	.estimator	mean	n	std_err	.config
	<int>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	1	rmse	standard	0.119	10	0.00417	Preprocessor1_Model01
2	2	rmse	standard	0.109	10	0.00393	Preprocessor1_Model02
3	3	rmse	standard	0.103	10	0.00371	Preprocessor1_Model03
4	4	rmse	standard	0.100	10	0.00360	Preprocessor1_Model04
5	5	rmse	standard	0.0987	10	0.00352	Preprocessor1_Model05
6	6	rmse	standard	0.0975	10	0.00347	Preprocessor1_Model06
7	7	rmse	standard	0.0969	10	0.00344	Preprocessor1_Model07
8	8	rmse	standard	0.0965	10	0.00342	Preprocessor1_Model08
9	9	rmse	standard	0.0963	10	0.00338	Preprocessor1_Model09
10	10	rmse	standard	0.0962	10	0.00333	Preprocessor1_Model10
11	11	rmse	standard	0.0962	10	0.00329	Preprocessor1_Model11
12	12	rmse	standard	0.0963	10	0.00328	Preprocessor1_Model12
13	13	rmse	standard	0.0965	10	0.00327	Preprocessor1_Model13

Plotting the Results

```
1 autoplot(knn_tune)
```



Selecting the Best Parameters

```
1 knn_best <- select_best(knn_tune)
2 knn_best
```

```
# A tibble: 1 × 2
  K .config
<int> <chr>
1    10 Preprocessor1_Model10
```

Getting the Best Model for Prediction

```
1 knn_final <-  
2   knn_wflow %>%  
3   finalize_workflow(knn_best)  
4  
5 knn_final
```

```
== Workflow ==  
Preprocessor: Formula  
Model: nearest_neighbor()  
  
— Preprocessor —  
Sale_Price ~ Longitude + Latitude  
  
— Model —  
K-Nearest Neighbor Model Specification (regression)  
  
Main Arguments:  
  neighbors = 10  
  
Computational engine: kknn
```

Note: Once we have this “final” model, we would still need to fit it with the **entire** training data set in order to then use it for prediction.

Other Tuning Parameters

- Boosting is an ensemble method that combines a series of base models, each of which is created sequentially and depends on the previous models
 - The number of boosting iterations is a tuning parameter
- In single-layer artificial neural network, the predictors are combined using two or more hidden units. The hidden units are linear combinations of the predictors that are captured in an *activation function* (typically a nonlinear function, such as a sigmoid).
 - The number of hidden units and the type of activation are tuning parameters.
- Modern gradient descent methods are improved by finding the right optimization parameters.
 - Examples of such hyperparameters are learning rates, momentum, and the number of optimization iterations/epochs.
 - Neural networks and some ensemble models use gradient descent to estimate the model parameters.

Tuning Preprocessing Steps

- In principal component analysis, the predictors are replaced with new, artificial features that have better properties related to collinearity.
 - The number of extracted components can be tuned.
- Imputation methods estimate missing predictor values using the complete values of one or more predictors. One effective imputation tool uses K -nearest neighbors of the complete columns to predict the missing value.
 - The number of neighbors can be tuned.

Tuning Structural Parameters

- In binary regression, the logit link is commonly used (i.e., logistic regression). Other link functions, such as the probit and complementary log-log, are also available and can be tuned.
- Non-Bayesian longitudinal and repeated measures models require a specification for the covariance or correlation structure of the data. Options include compound symmetric (a.k.a. exchangeable), autoregressive, Toeplitz, and others, which can be tuned.

Two general strategies for optimization

Tuning parameter optimization usually falls into one of two categories: *grid search* and *iterative search*.

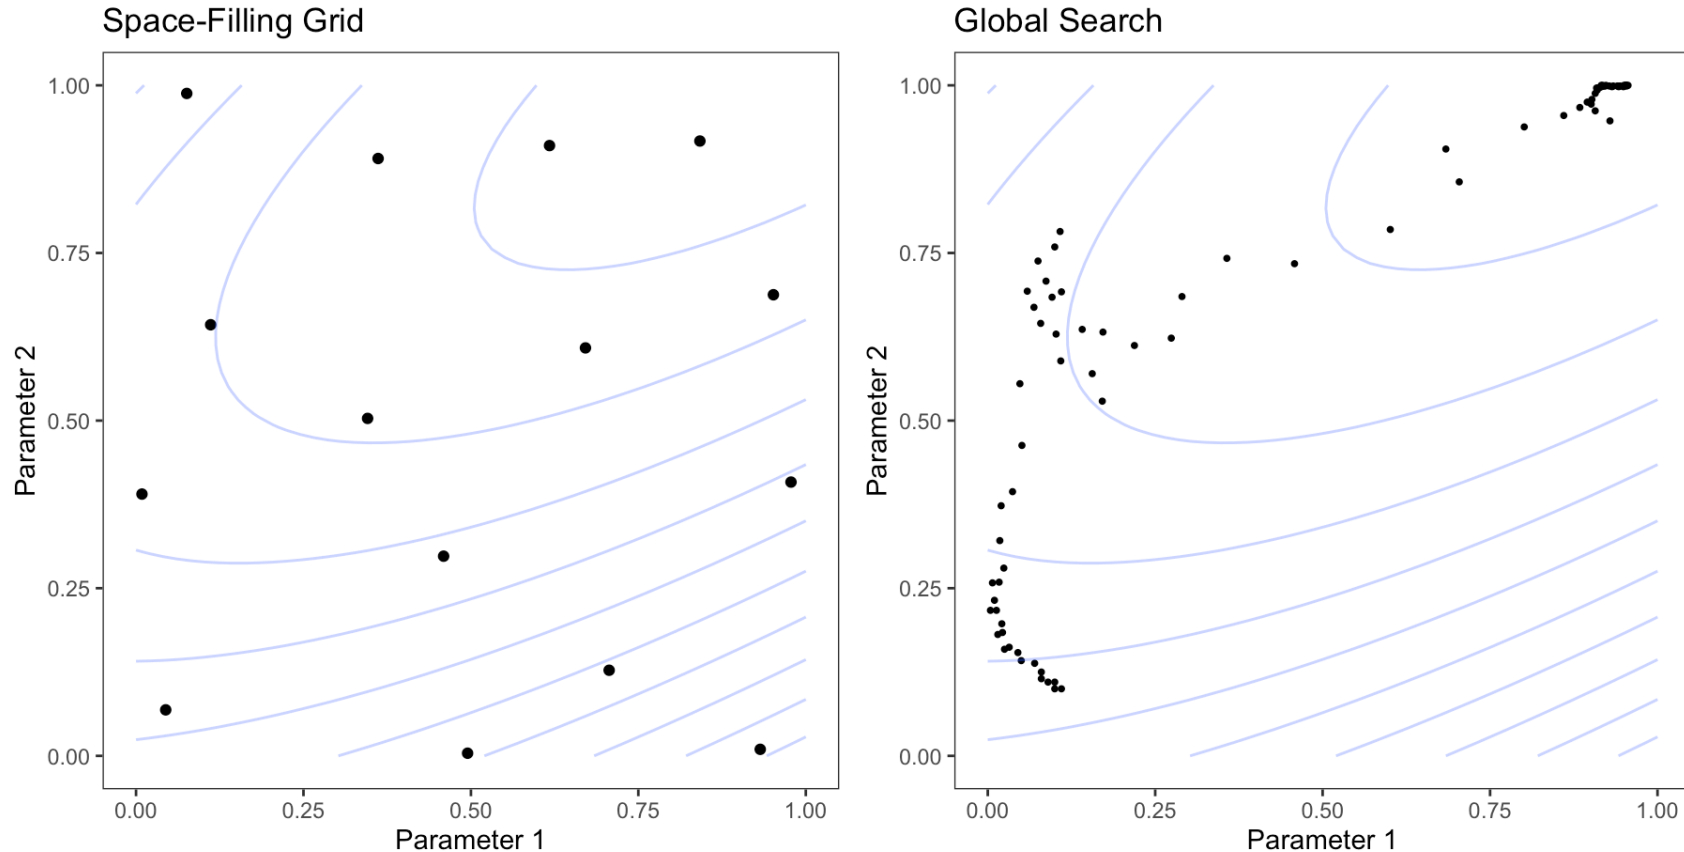
Grid Search

- *Grid search* is when we predefine a set of parameter values to evaluate
- The main choices involved in grid search are how to make the grid and how many parameter combinations to evaluate
- Grid search is often judged as inefficient since the number of grid points required to cover the parameter space can become unmanageable with the curse of dimensionality
- Lots of times it gets the job done

Iterative Search

- *Iterative search* or sequential search is when we sequentially discover new parameter combinations based on previous results
- Almost any nonlinear optimization method is appropriate, although some are more efficient than others.
- In some cases, an initial set of results for one or more parameter combinations is required to start the optimization process.

Visualizing the Two Approaches



Examples of pre-defined grid tuning and an iterative search method. The lines represent contours of a performance metric; it is best in the upper-right-hand side of the plot.